

**KUBERNETES**

Kubernetes is a open-source platform for managing containerized workloads and services.

Kubernetes facilitates declarative configuration and automation.

Kubernetes originates from Greek term meaning helmsman or pilot.

Google open-sourced the Kubernetes project in 2014.

Containers are a good way to bundle and run your applications.

In a production environment, containers that run the applications should be always running without downtime.

For example, if a container goes down, another container needs to start.

That is Kubernetes orchestration.

Kubernetes (K8s) is an open source container orchestration tool originally developed and maintained google

Originally developed by google and now maintained by Cloud Native Computing Foundation (CNCF)

Kubernetes orchestration tool can handle:

Automated binpacking

Provision and deployment of containers

Scaling up or Scaling down containers depending on the load.

Movement of containers from one host to another host depending on the resources.

Loadbalancing of containers based on the request by users.

Healing of containers

Rollouts and Rollbacks

## **Container orchestration solutions**

Docker Swarm

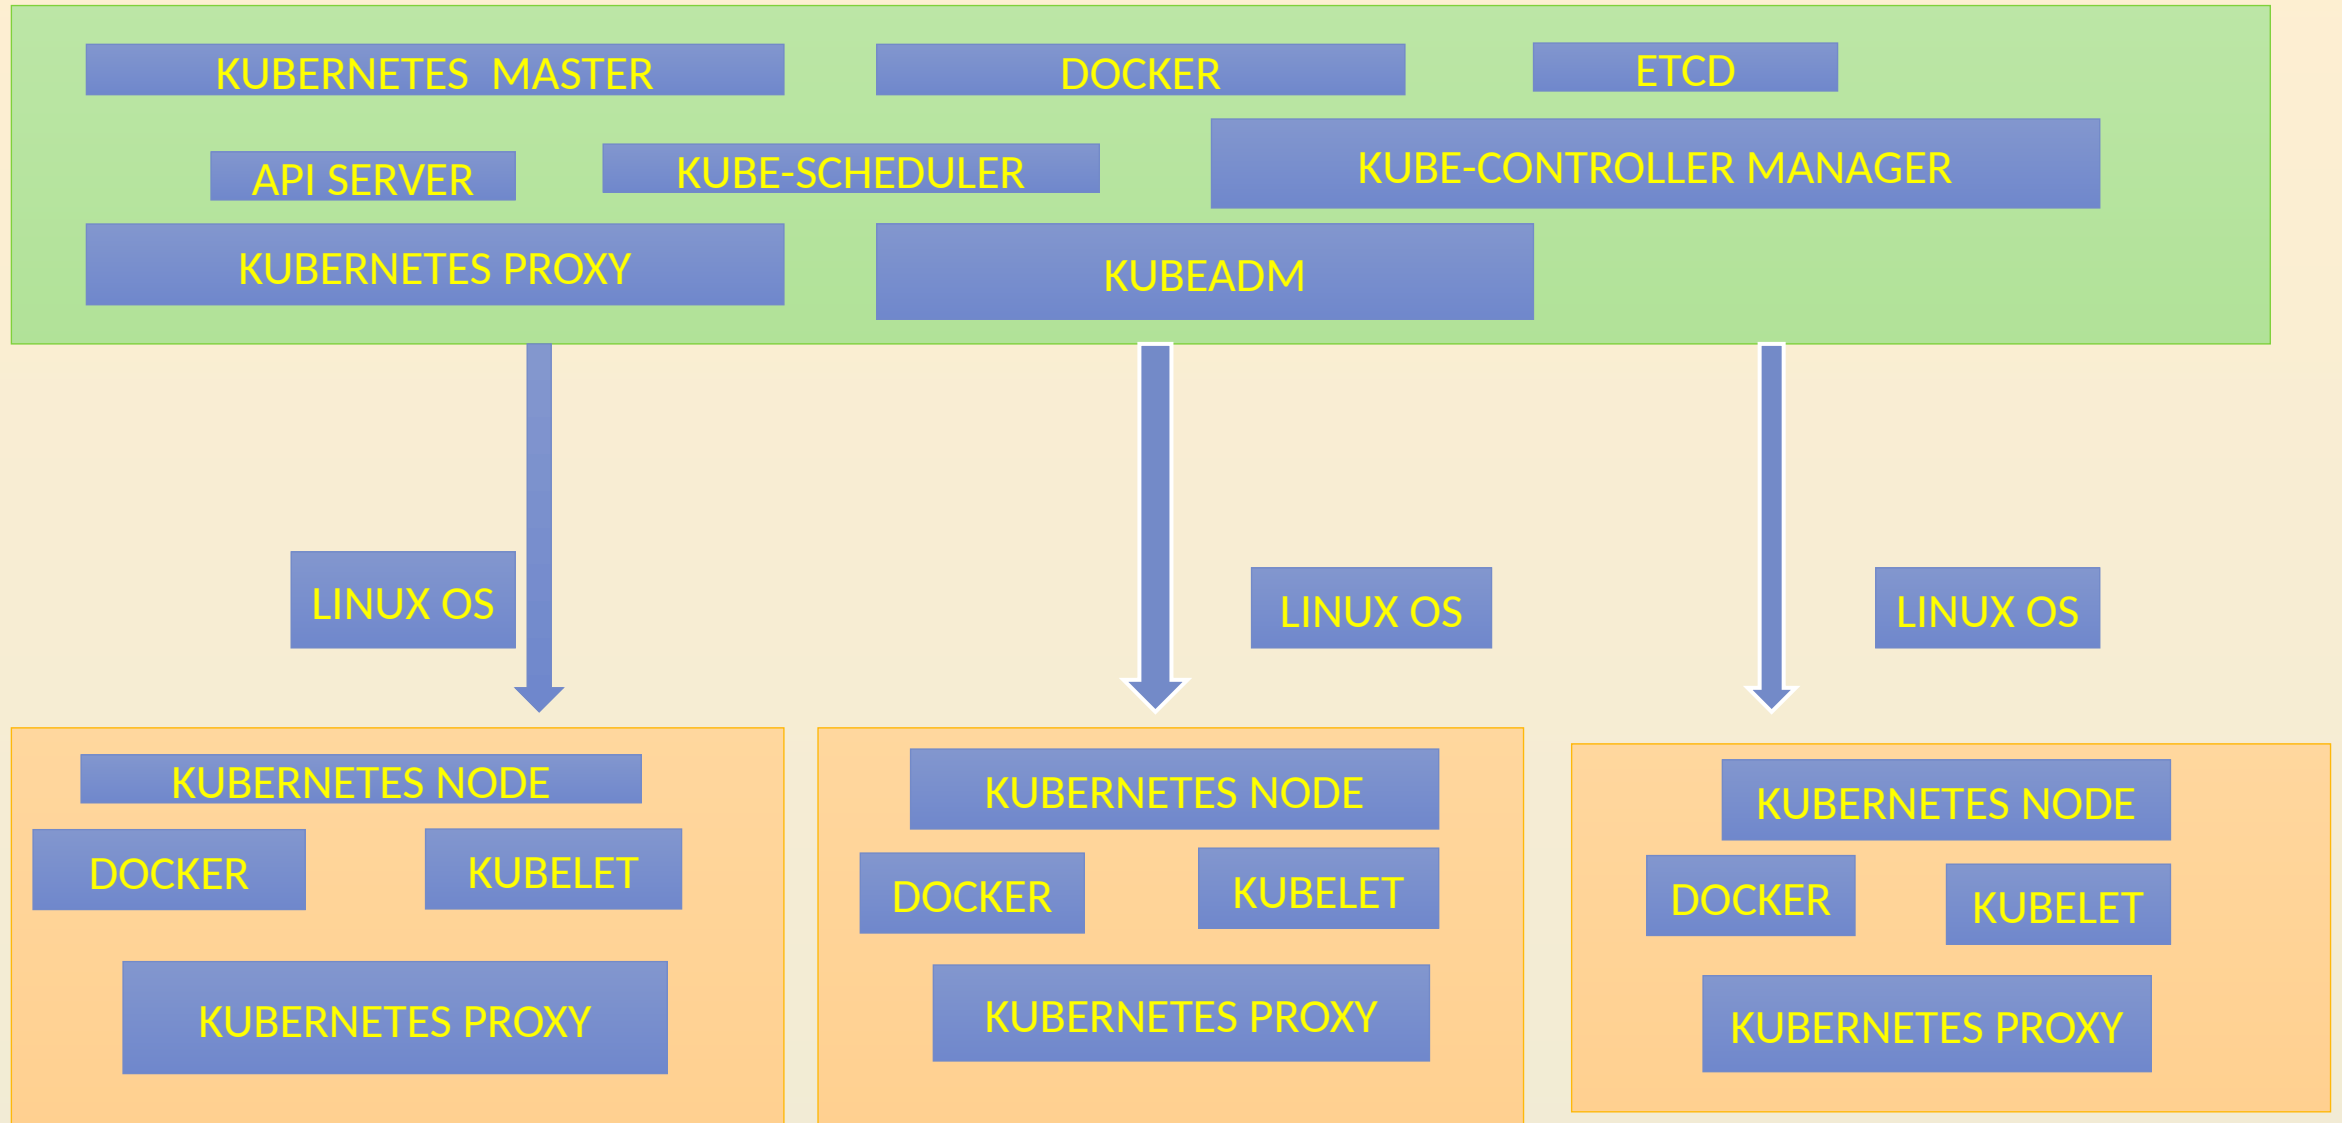
Kubernetes

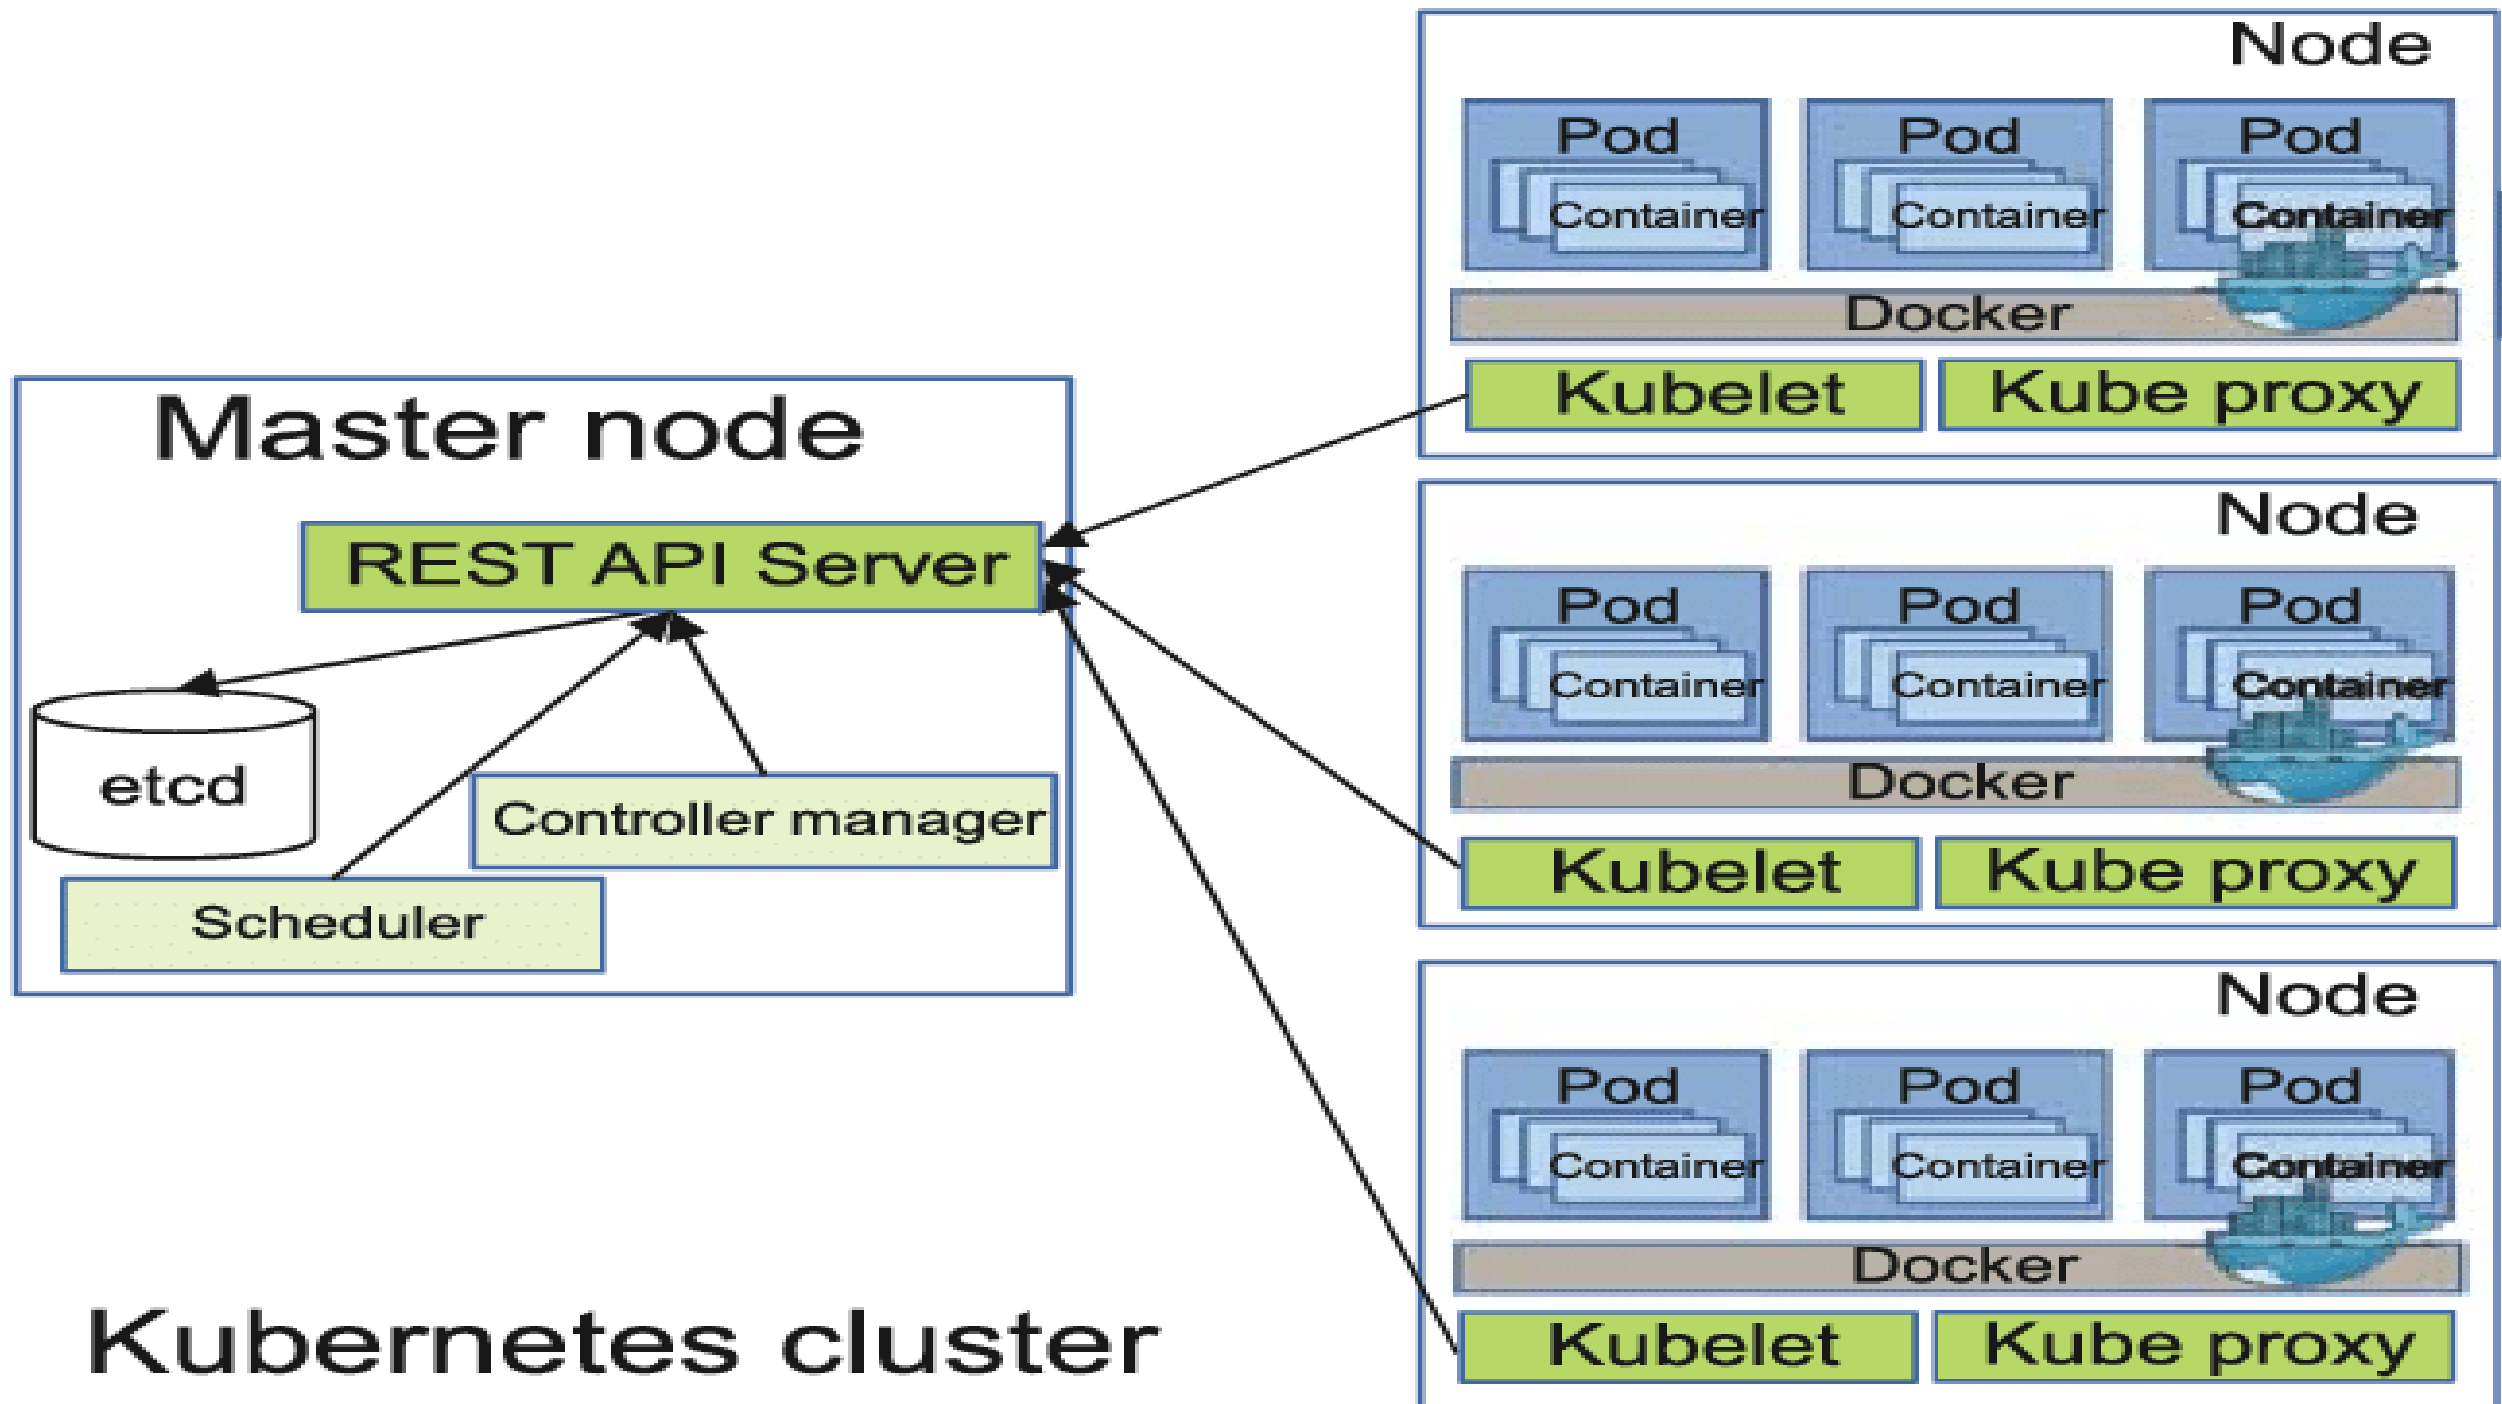
Apache Mesos

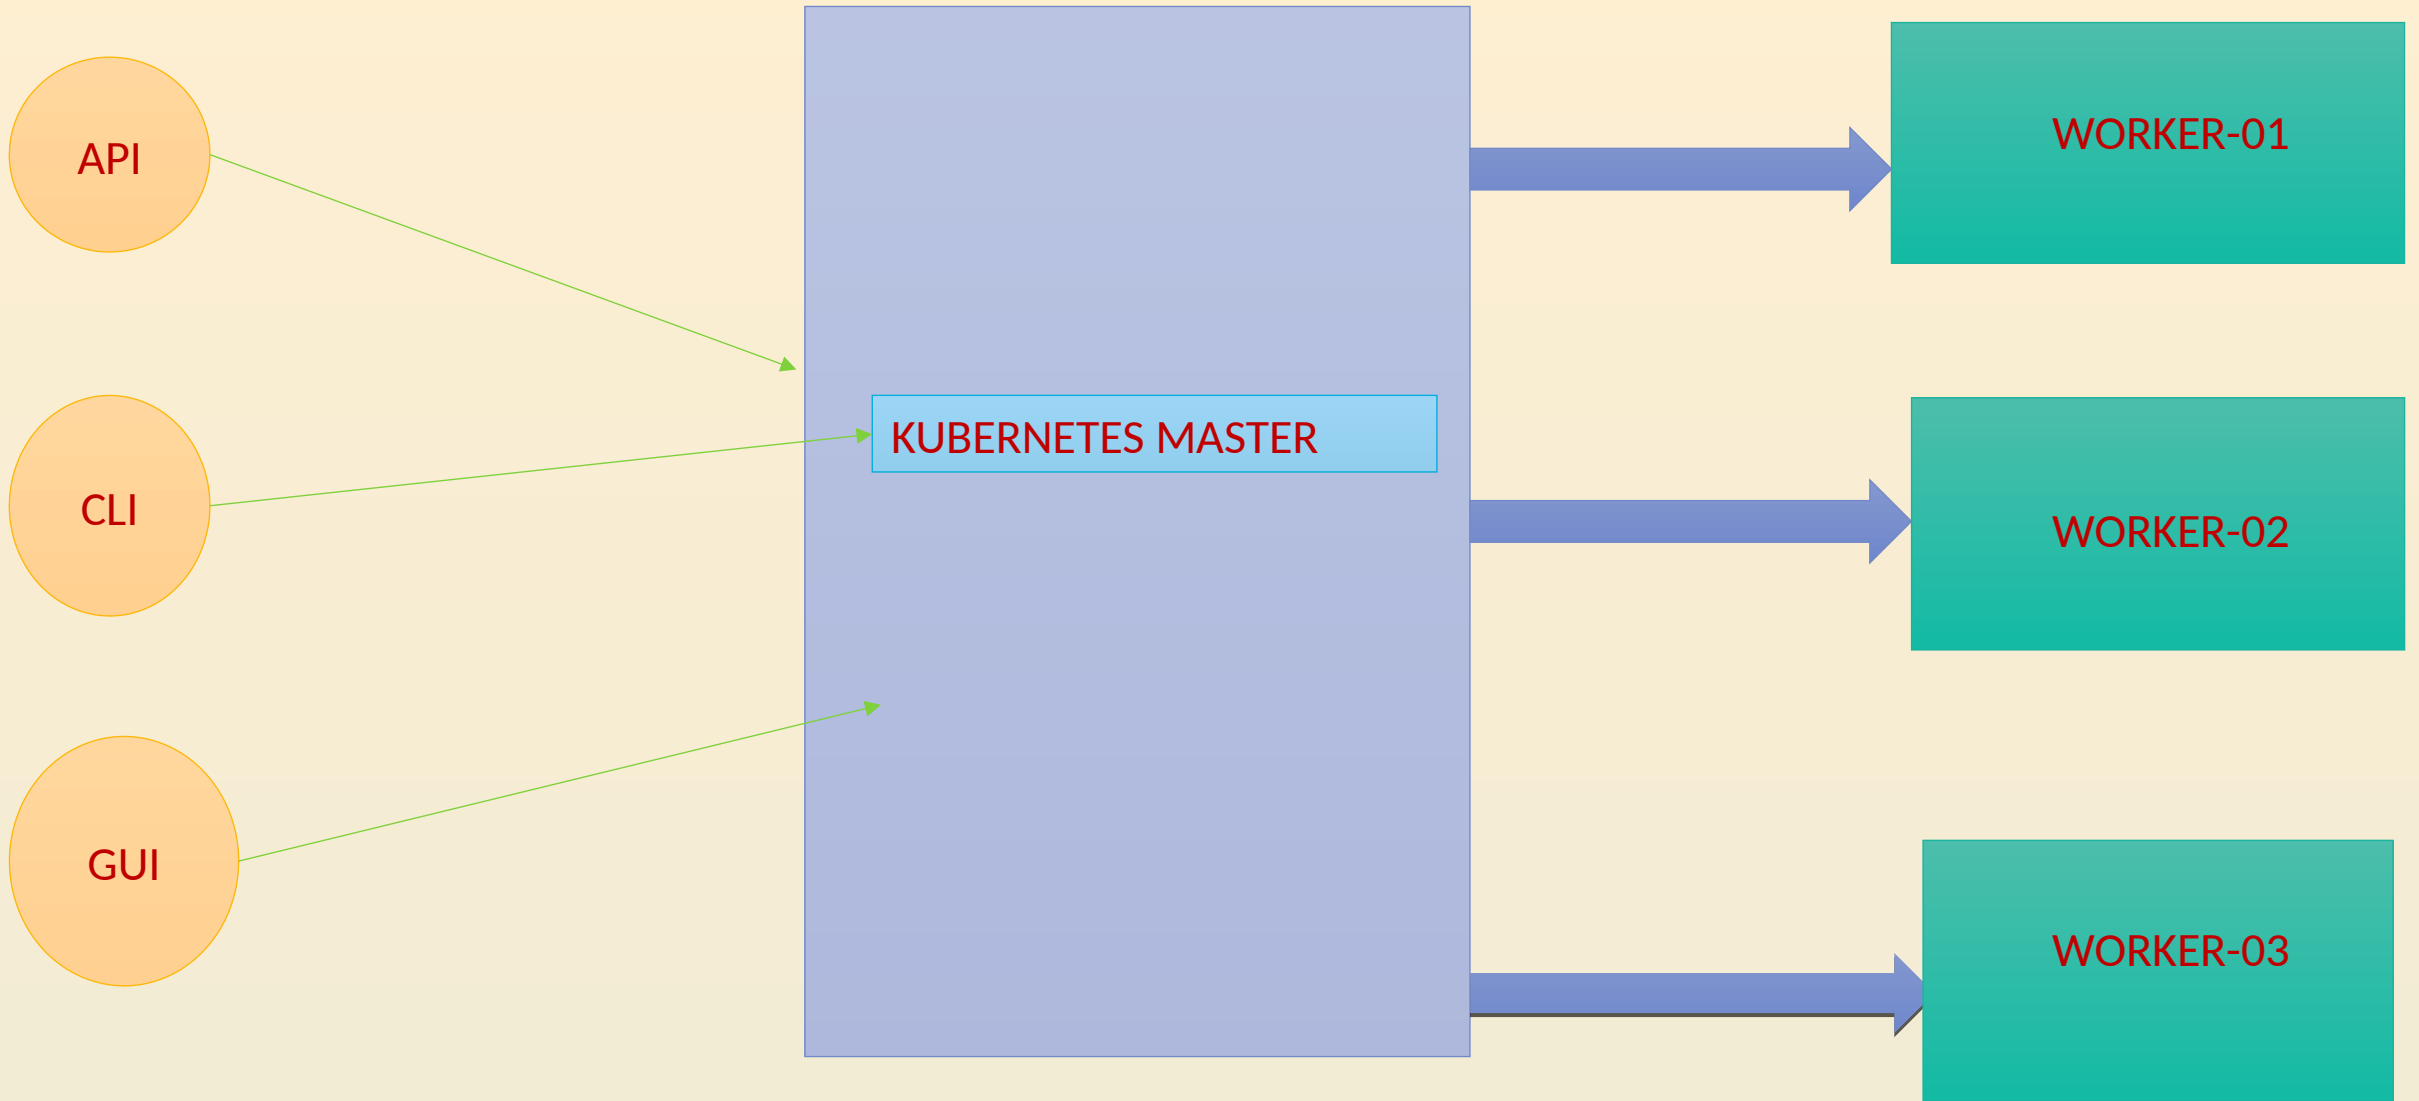
Elastic Kubernetes Service (AWS EKS)

Google Container Engine (GKE)

# Kubernetes - Cluster Architecture







## KUBERNETES INSTALLATION METHODS

Use managed Kubernetes Service from various providers like AWS, AZURE, GOOGLE, IBM, DigitalOcean and others.

Minikube ( Installation of kubernetes locally a single node cluster)

Install and Configure Kubernetes Manually (Hard Way)



# POD

*Pods* are the smallest deployable units of computing that you can create and manage in Kubernetes.

A *Pod* is a group of one or more containers.

**Pods that run a single container.** The "one-container-per-Pod" model is the most common Kubernetes use case.

in this case, you can think of a Pod as a wrapper around a single container.

Kubernetes manages Pods rather than managing the containers directly.

**Pods that run multiple containers that need to work together.**

A Pod can run co-located containers form a single cohesive unit of service

For example,

one container serving data stored in a shared volume to the public,

while a separate *sidecar* container refreshes or updates those files.

Each Pod is meant to run a single instance of a given application.

If you want to scale your application horizontally (to provide more overall resources by running more instances),

you should use multiple Pods, one for each instance.

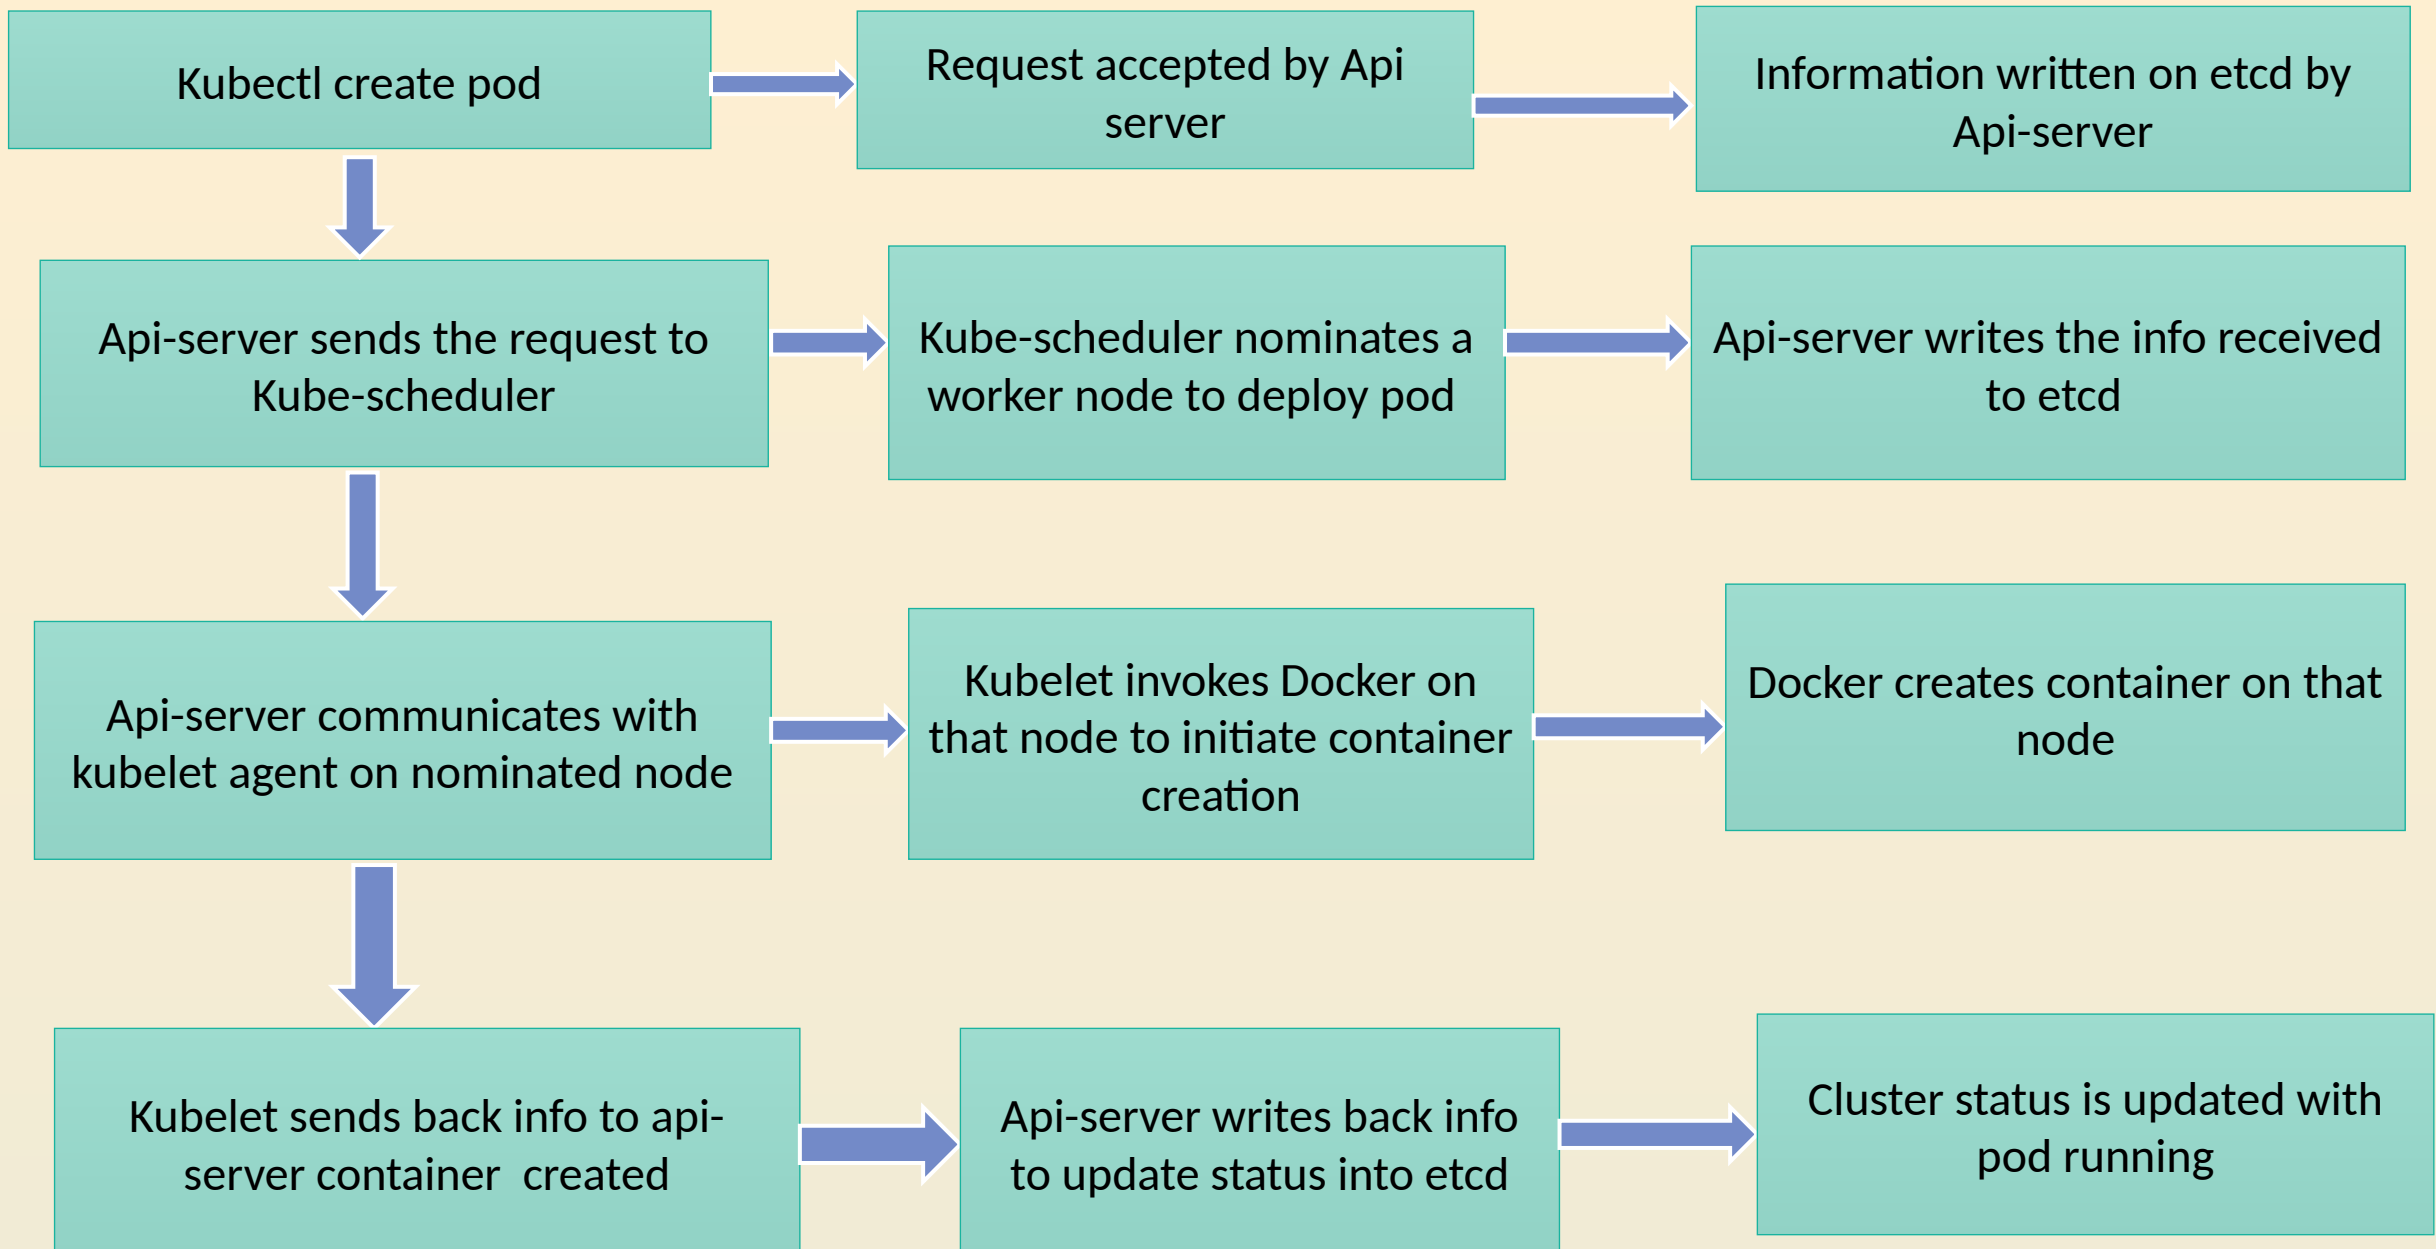
In Kubernetes, this is typically referred to as *replication*.

Replicated Pods are usually created and managed as a group by a **controller**.

A Pod is not a process, but an environment for running container(s).

A Pod persists until it is deleted.

## Flow of Deployment of a Pod in Kubernetes Cluster



# Replicaset

A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time.

It is often used to guarantee the availability of a specified number of identical Pods.

## How a ReplicaSet works




Number of replicas indicate how many Pods it should be maintaining,

A pod template specifying the data of new Pods it should create to meet the number of replicas criteria.


When a ReplicaSet needs to create new Pods, it uses its Pod template.

A ReplicaSet then fulfills its purpose by creating and deleting Pods as needed to reach the desired number.

# REPLICASET



# Replica Set



# DEPLOYMENT

Deployment manages replica sets which is also an upgraded version of the replication controller.

Deployment has capability to update the replica set and are also capable of rolling back to the previous version.

# Deployment

## REPLICASET

POD

CONTAINER


POD

CONTAINER

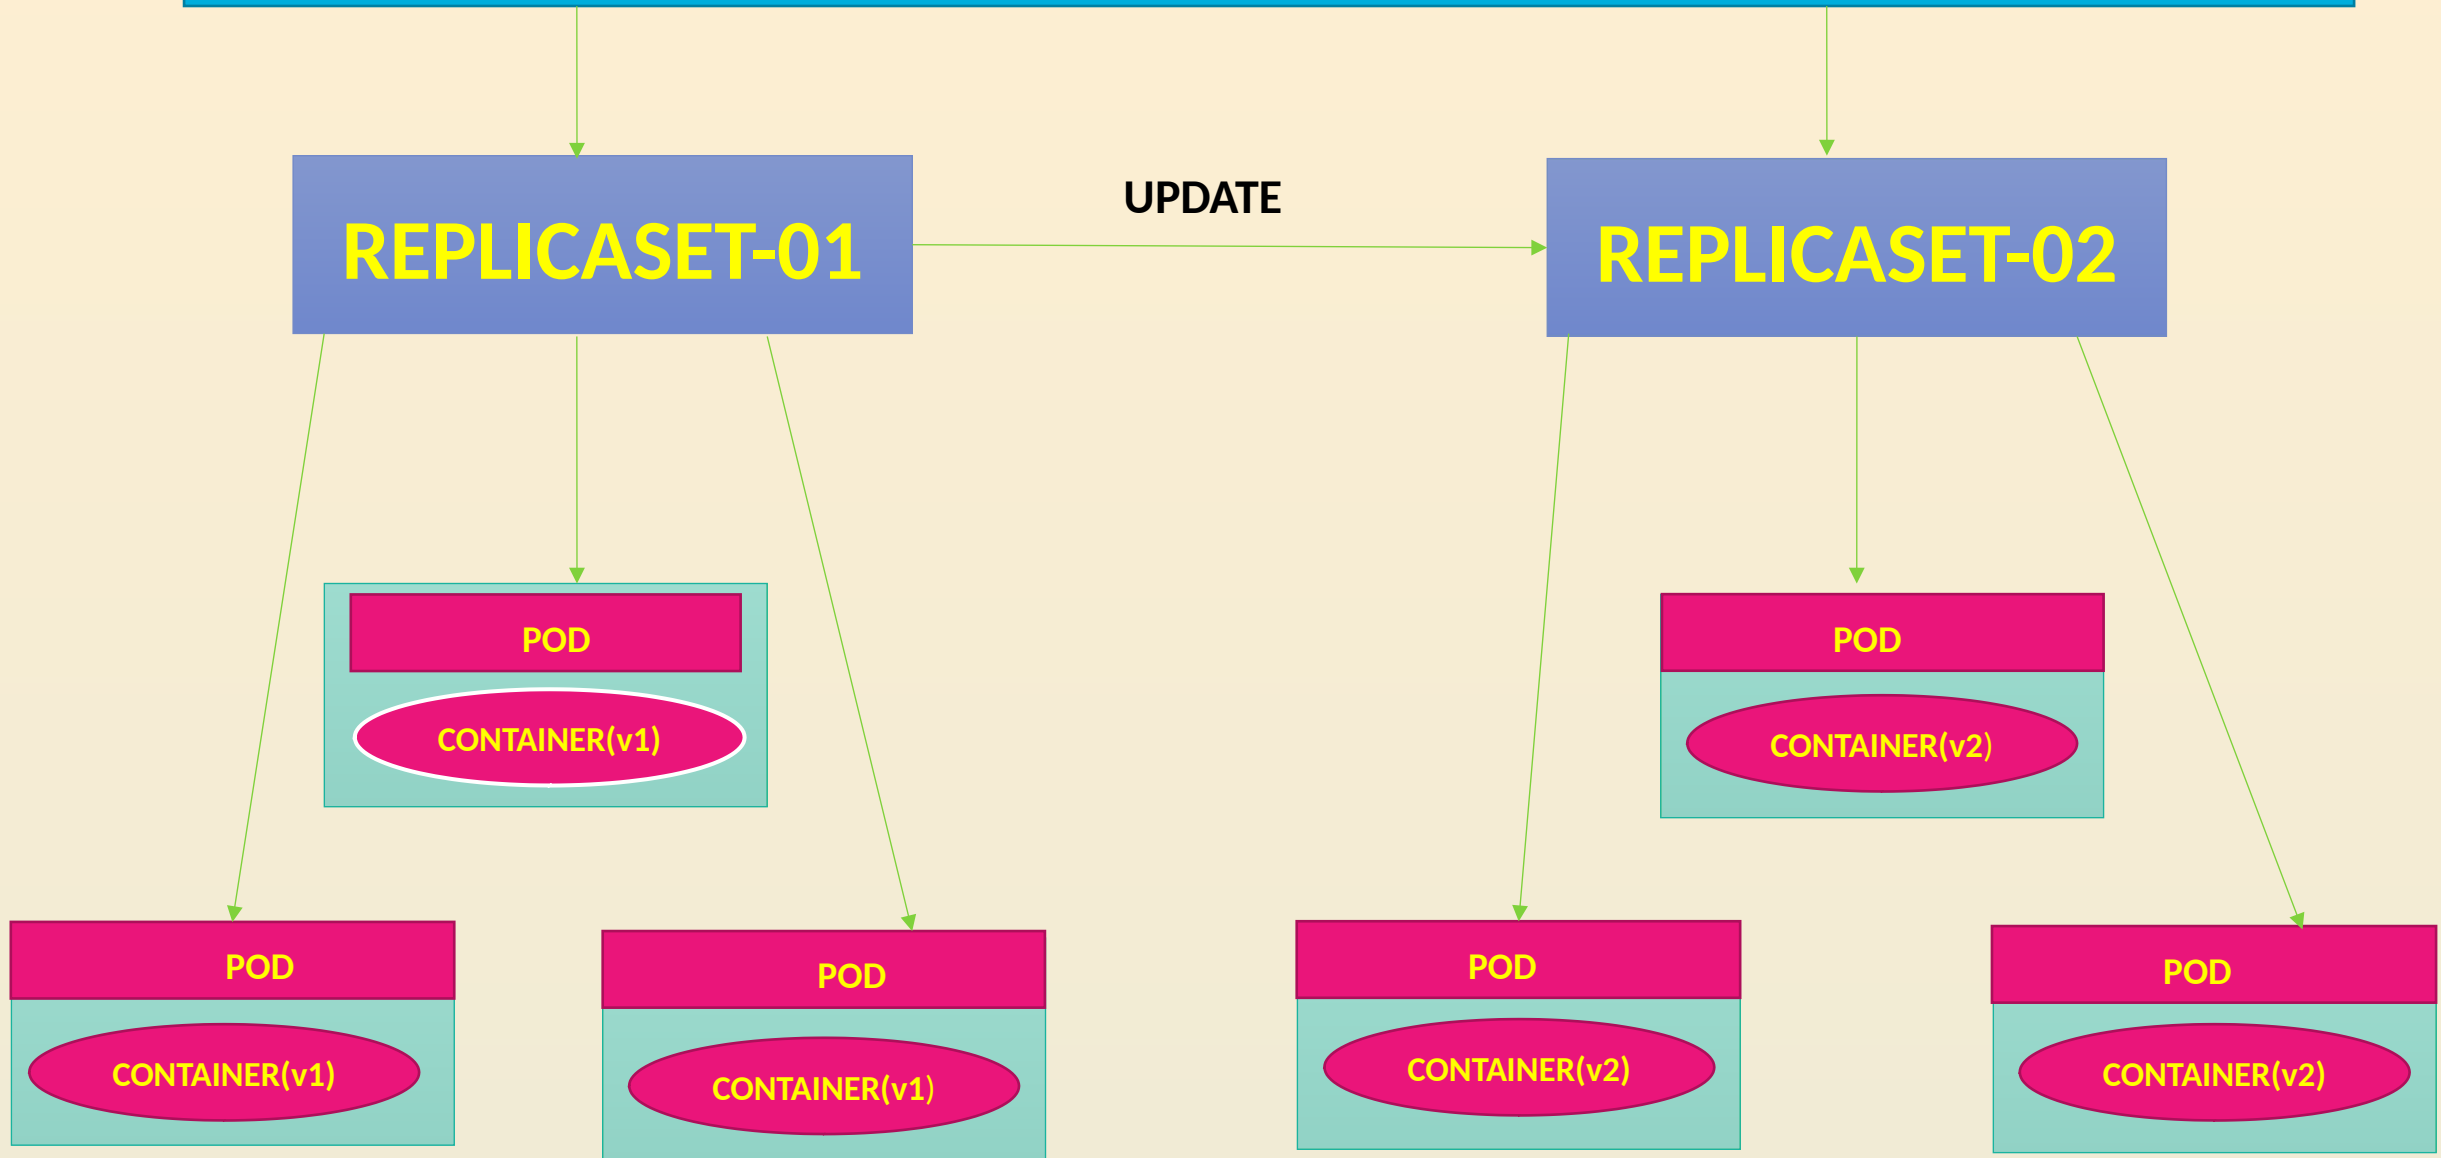
POD

CONTAINER





# Deployment



## Kubernetes service

Kubernetes **Pods** are created and destroyed to match the state of your cluster.

Pods are nonpermanent resources.

If you use a **Deployment** to run your app, it can create and destroy Pods dynamically.

Each Pod gets its own IP address.

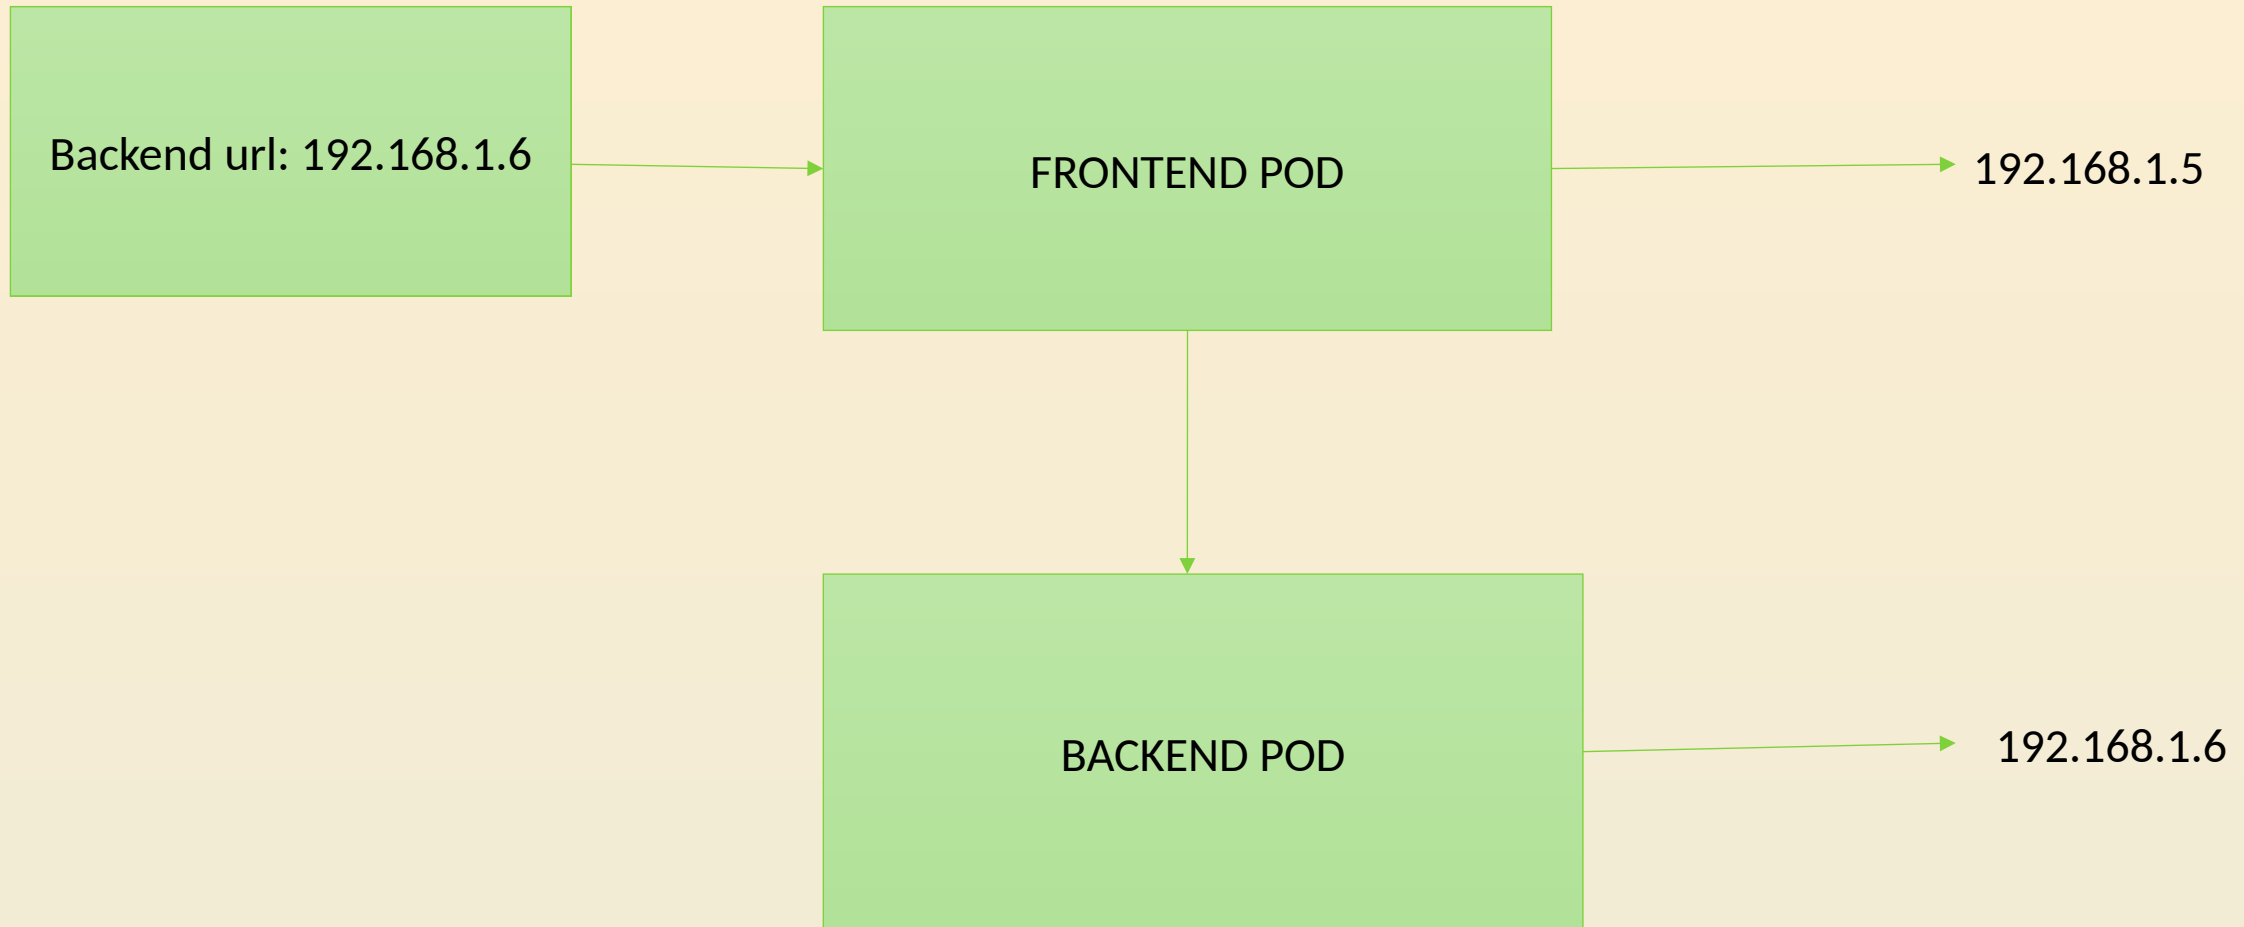
This leads to a problem:

If some set of Pods (call them "backends") provides functionality to other Pods (call them "frontends") inside your cluster,

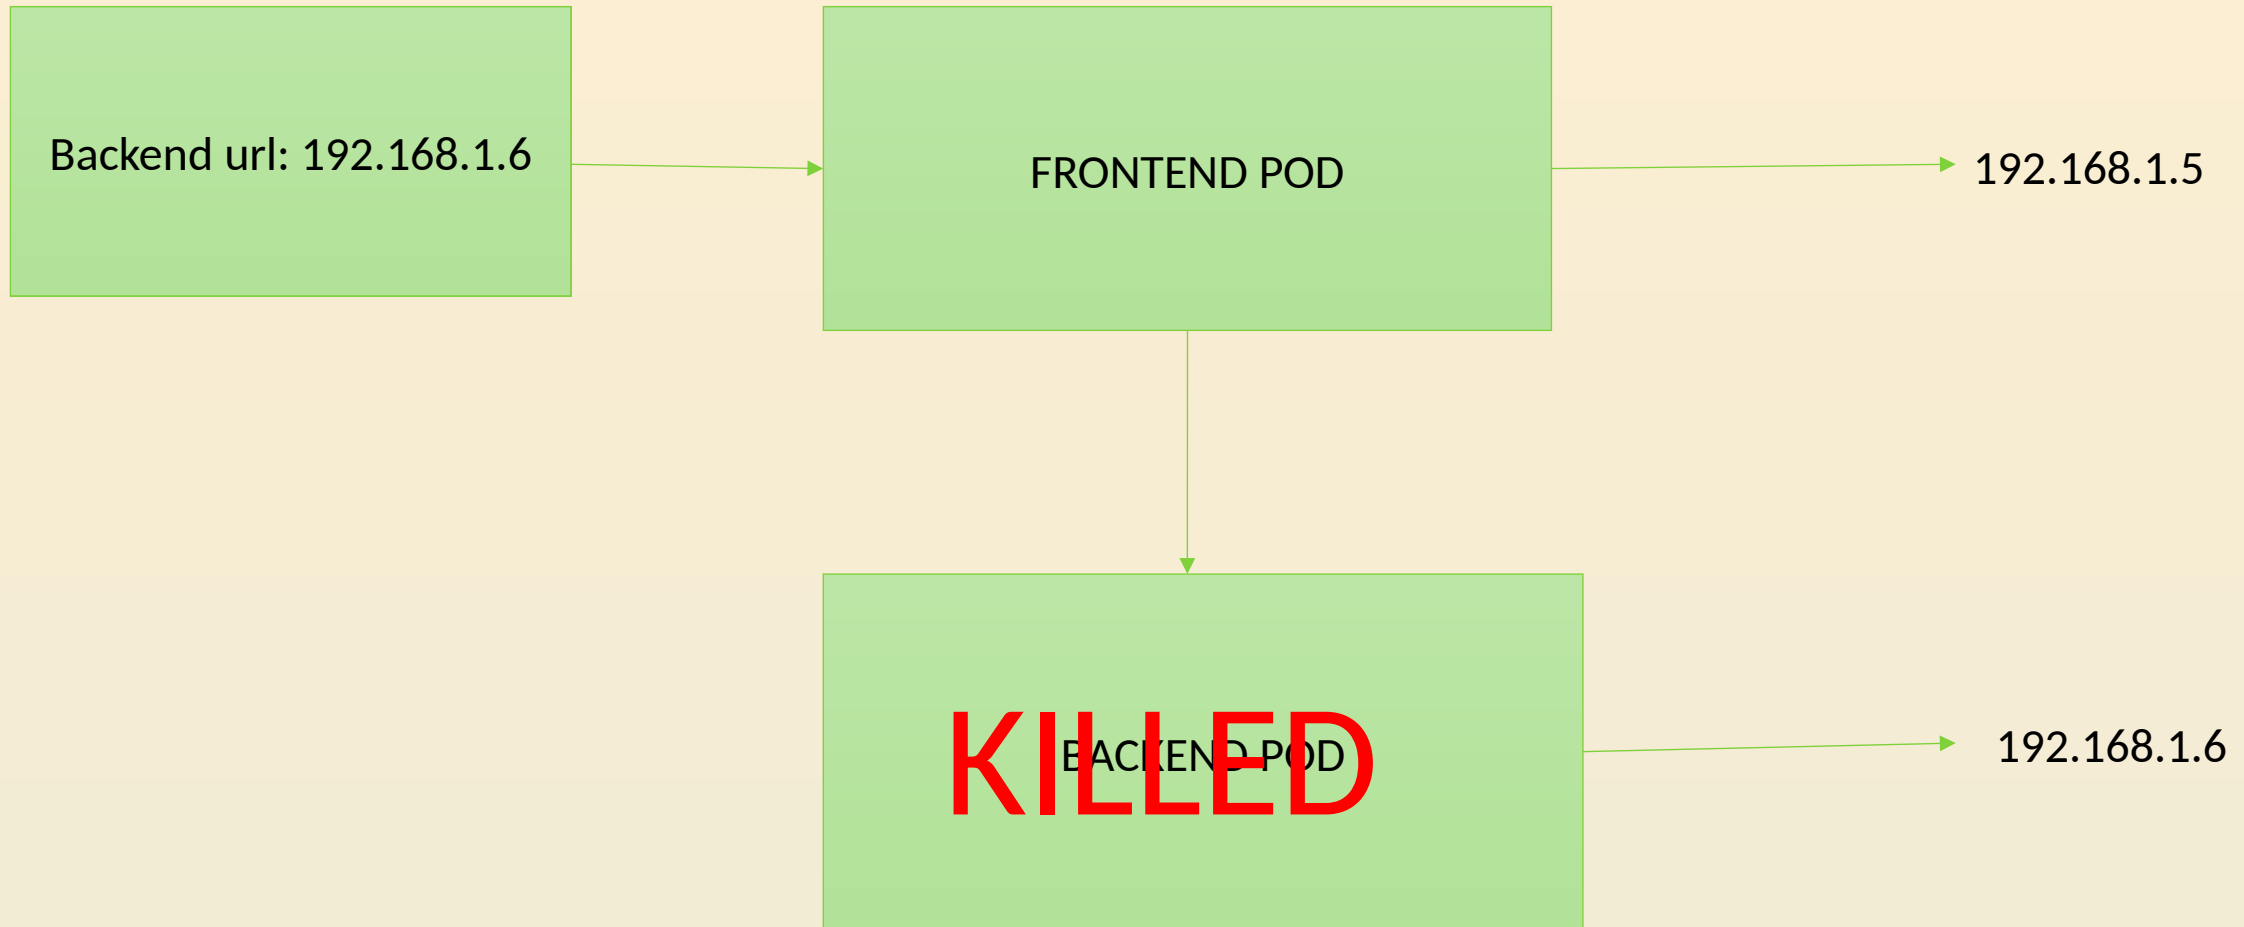
How do the frontends find out and keep track of which IP address to connect to?

## KUBERNETES SERVICE

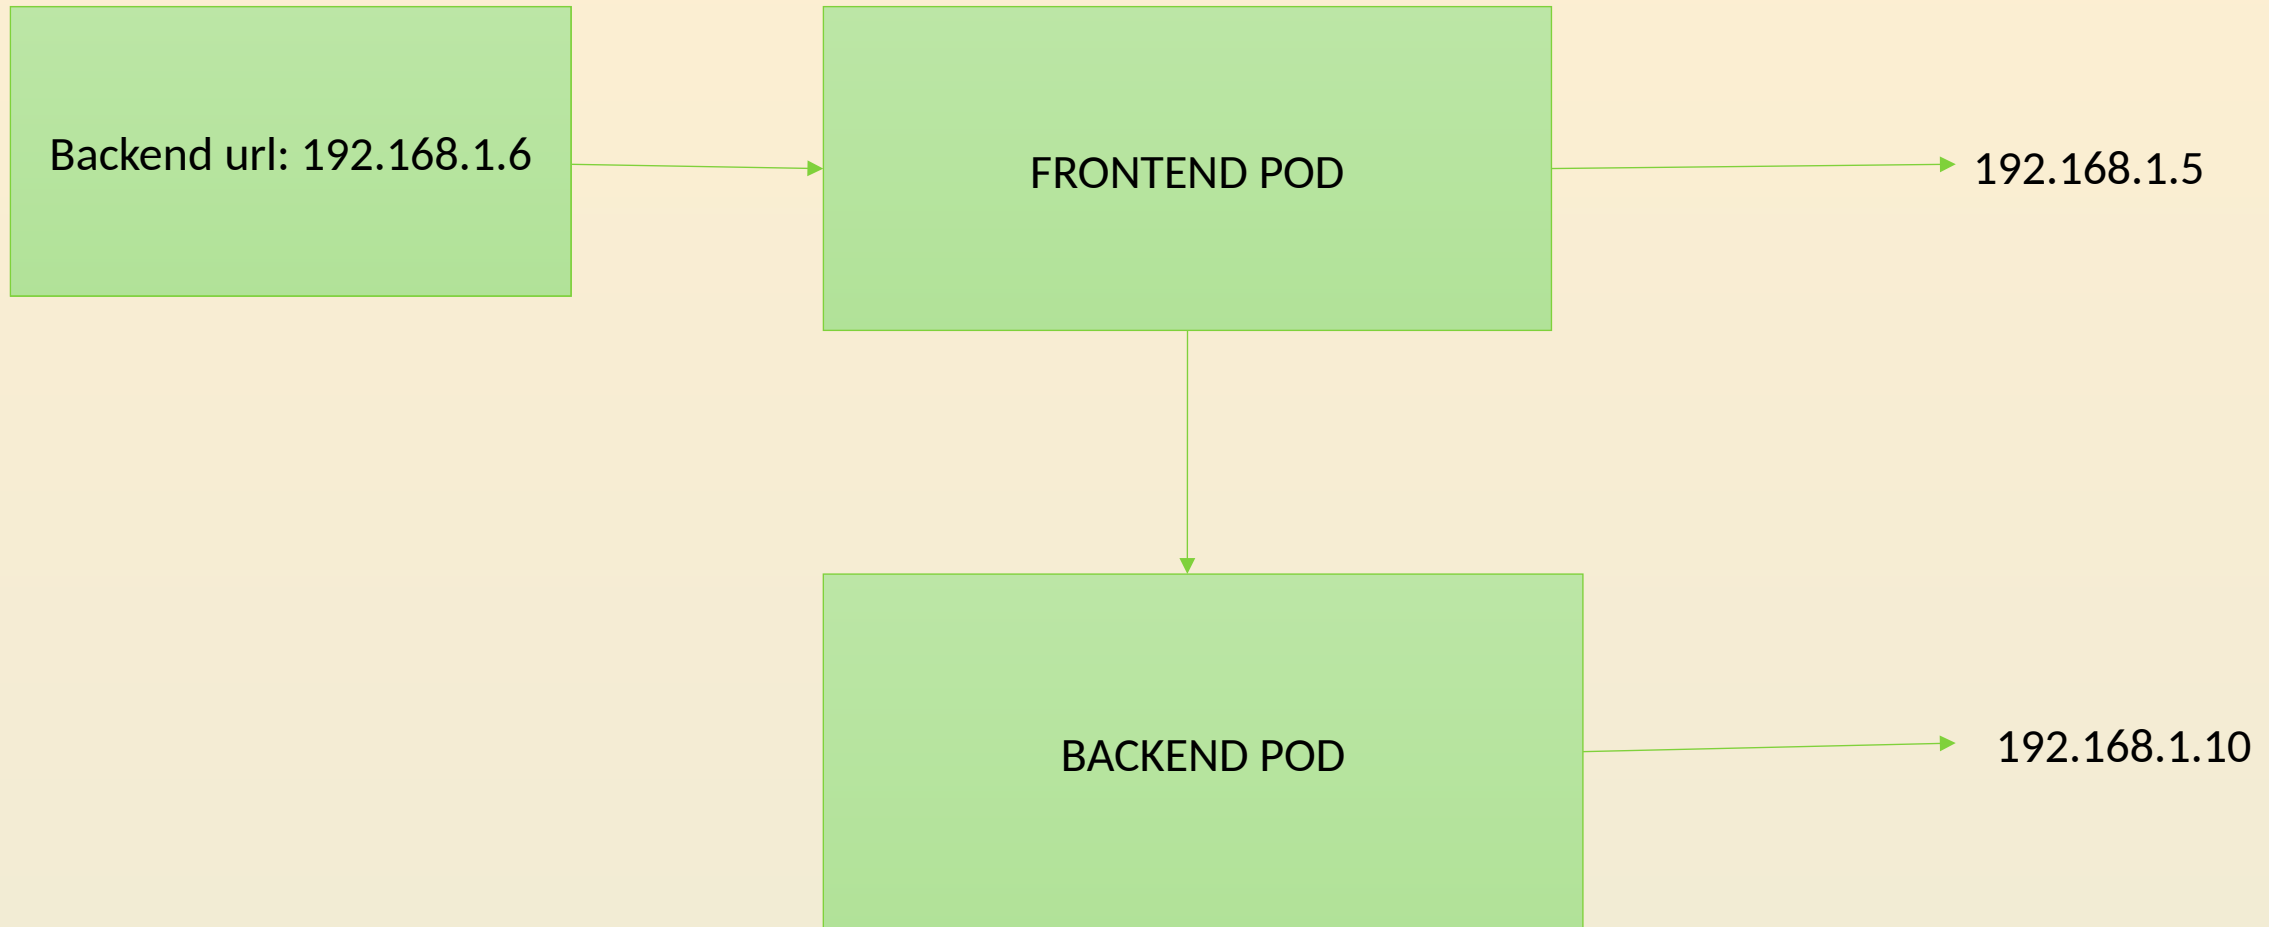
When you create a pod each pod has an unique private ip address



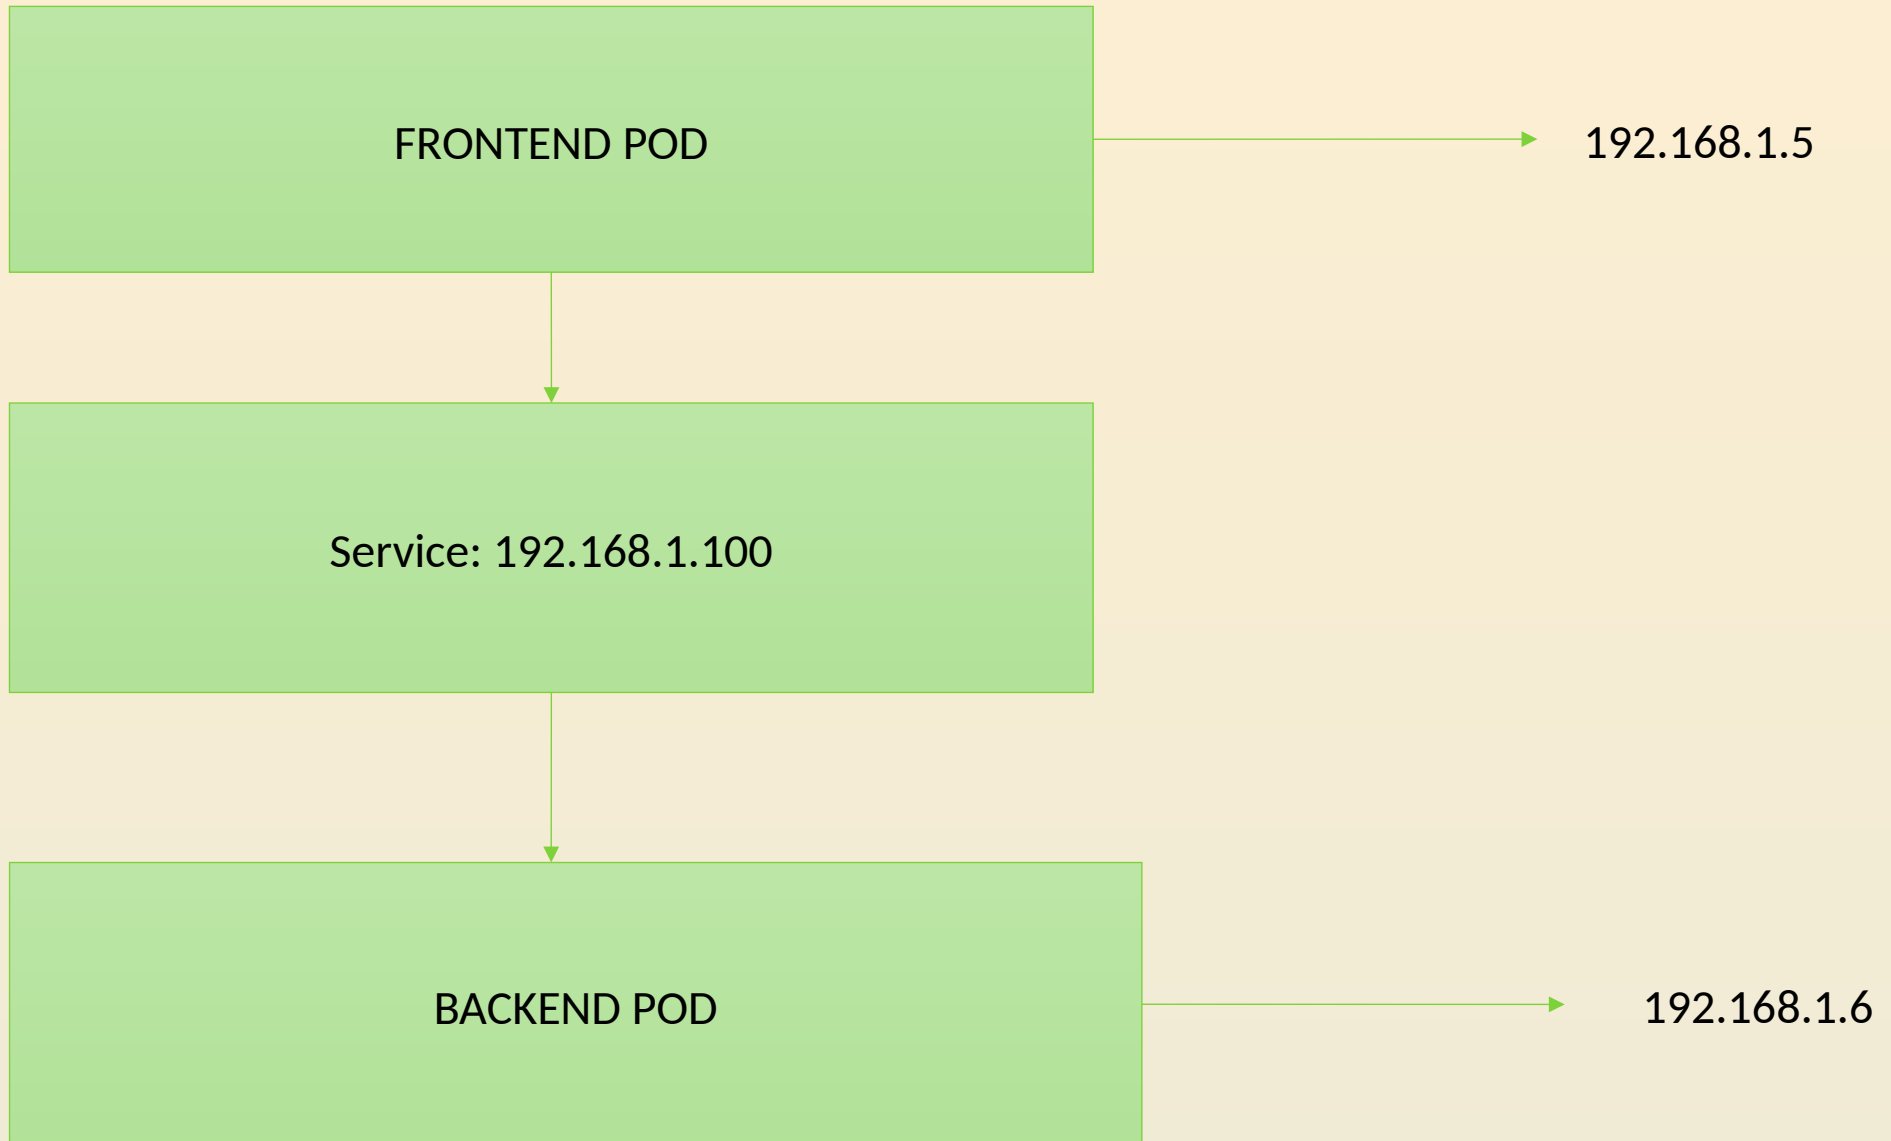
**Pods are Ephemeral, pods are recreated**



**New created pod has an different private ip address**



**You can create a service which has static ip and act as gateway for your frontend to talk to backend pod**



# Types of Services in Kubernetes

ClusterIP

NodePort

LoadBalancer



## ClusterIP

Whenever the service type is ClusterIP, an internal ip of the cluster is assigned to the service.

Since an internal ip is assigned, it can only be reachable within the cluster.

This is the default service type.

## Front-End Pods

Web App

Web App

Web App

## Backend Pods

Backend Pod

Backend Pod

Backend Pod

## Redis Pods

Redis Pod

Redis Pod

Redis Pod

Backend ( Cluster IP)

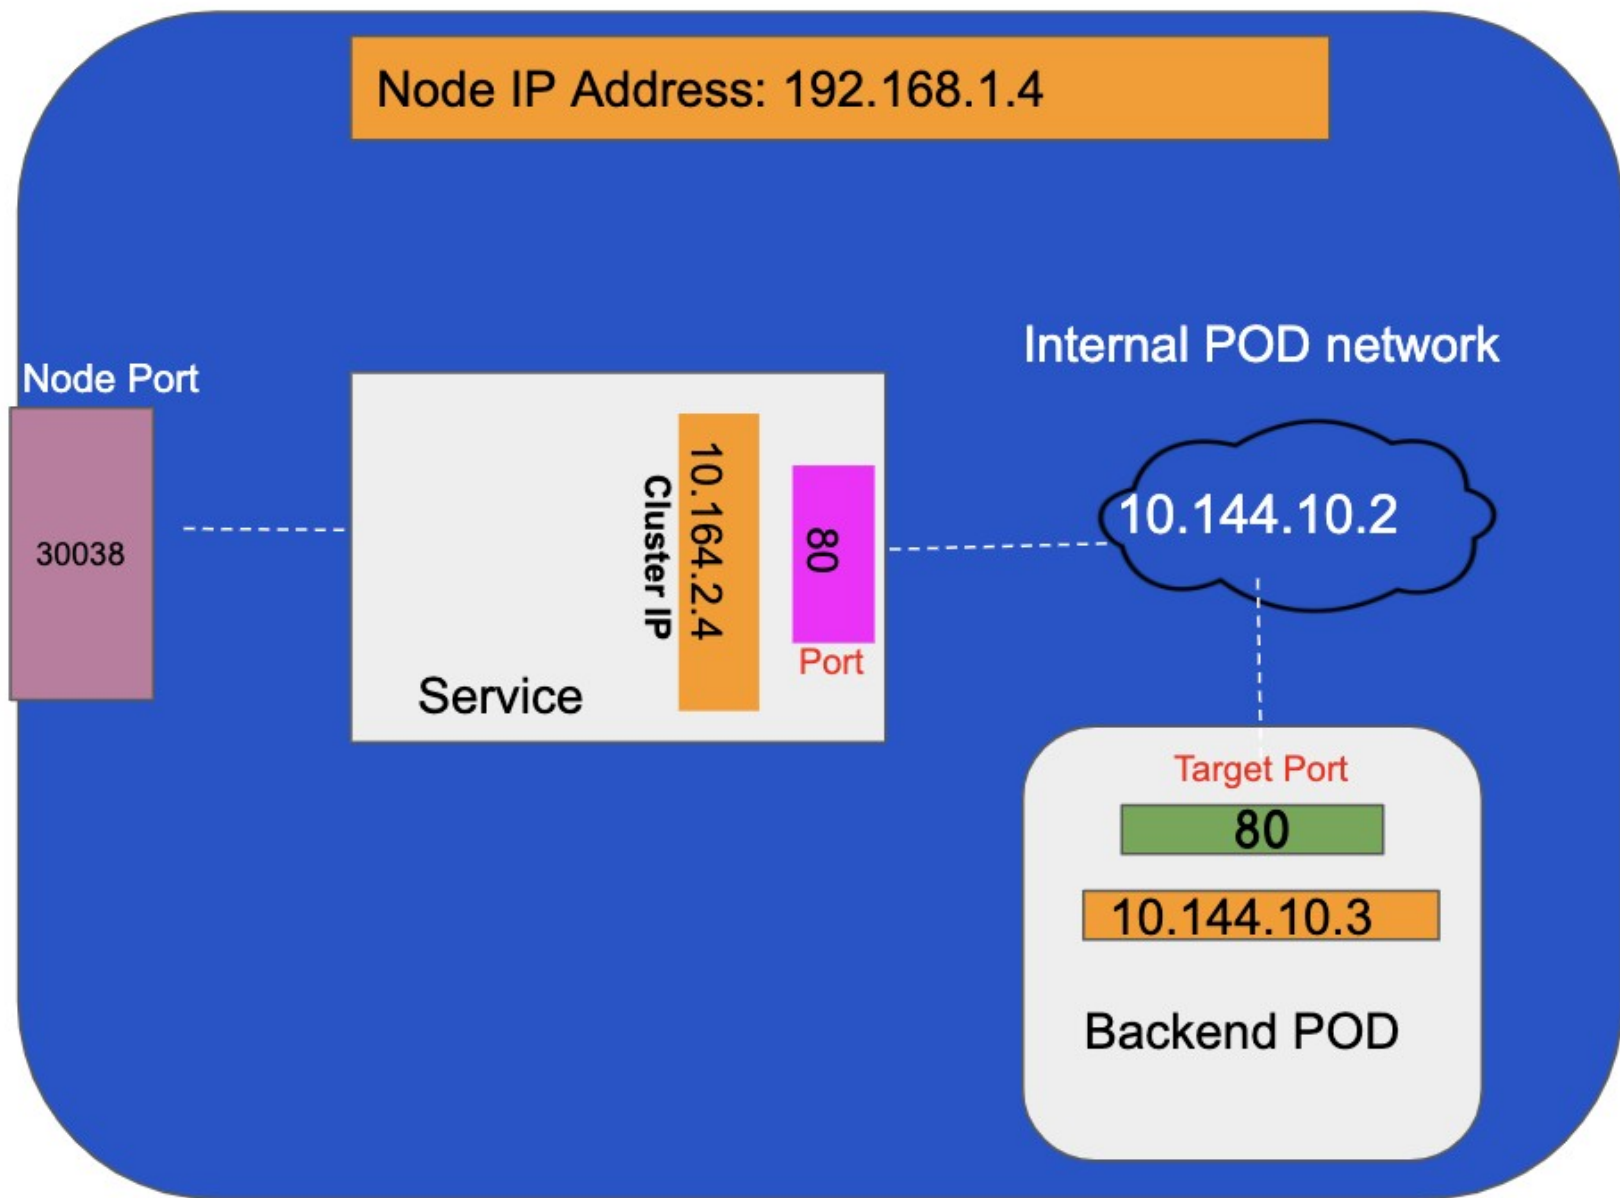
Redis ( Cluster IP)

## **NODEPORT**

Nodeport service opens the port on the nodes.

If service type is Nodeport, then kubernetes will allocate a port (default- 30000-32767) on all the nodes in the cluster.

Each node will proxy that port into your service.



NODE

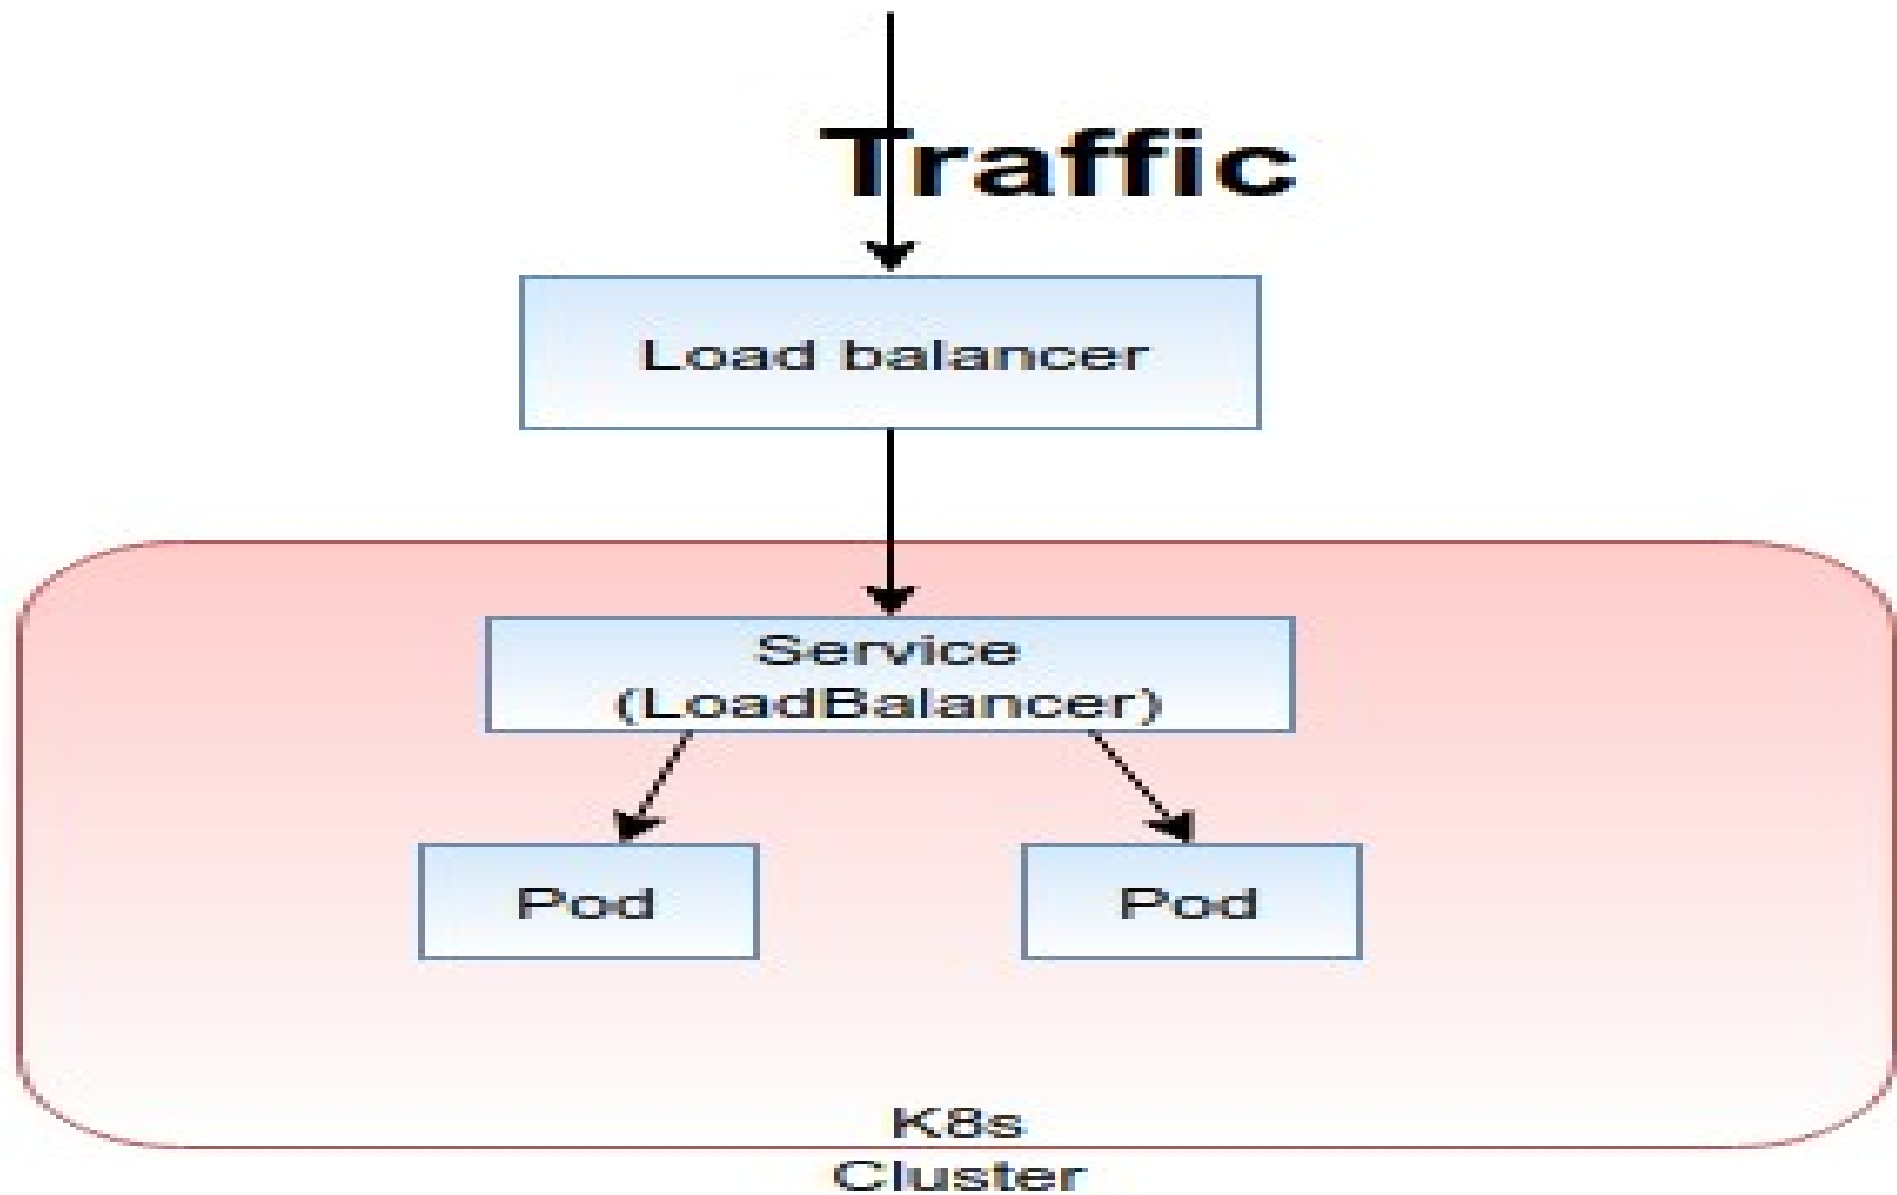
## **LOADBALANCER**

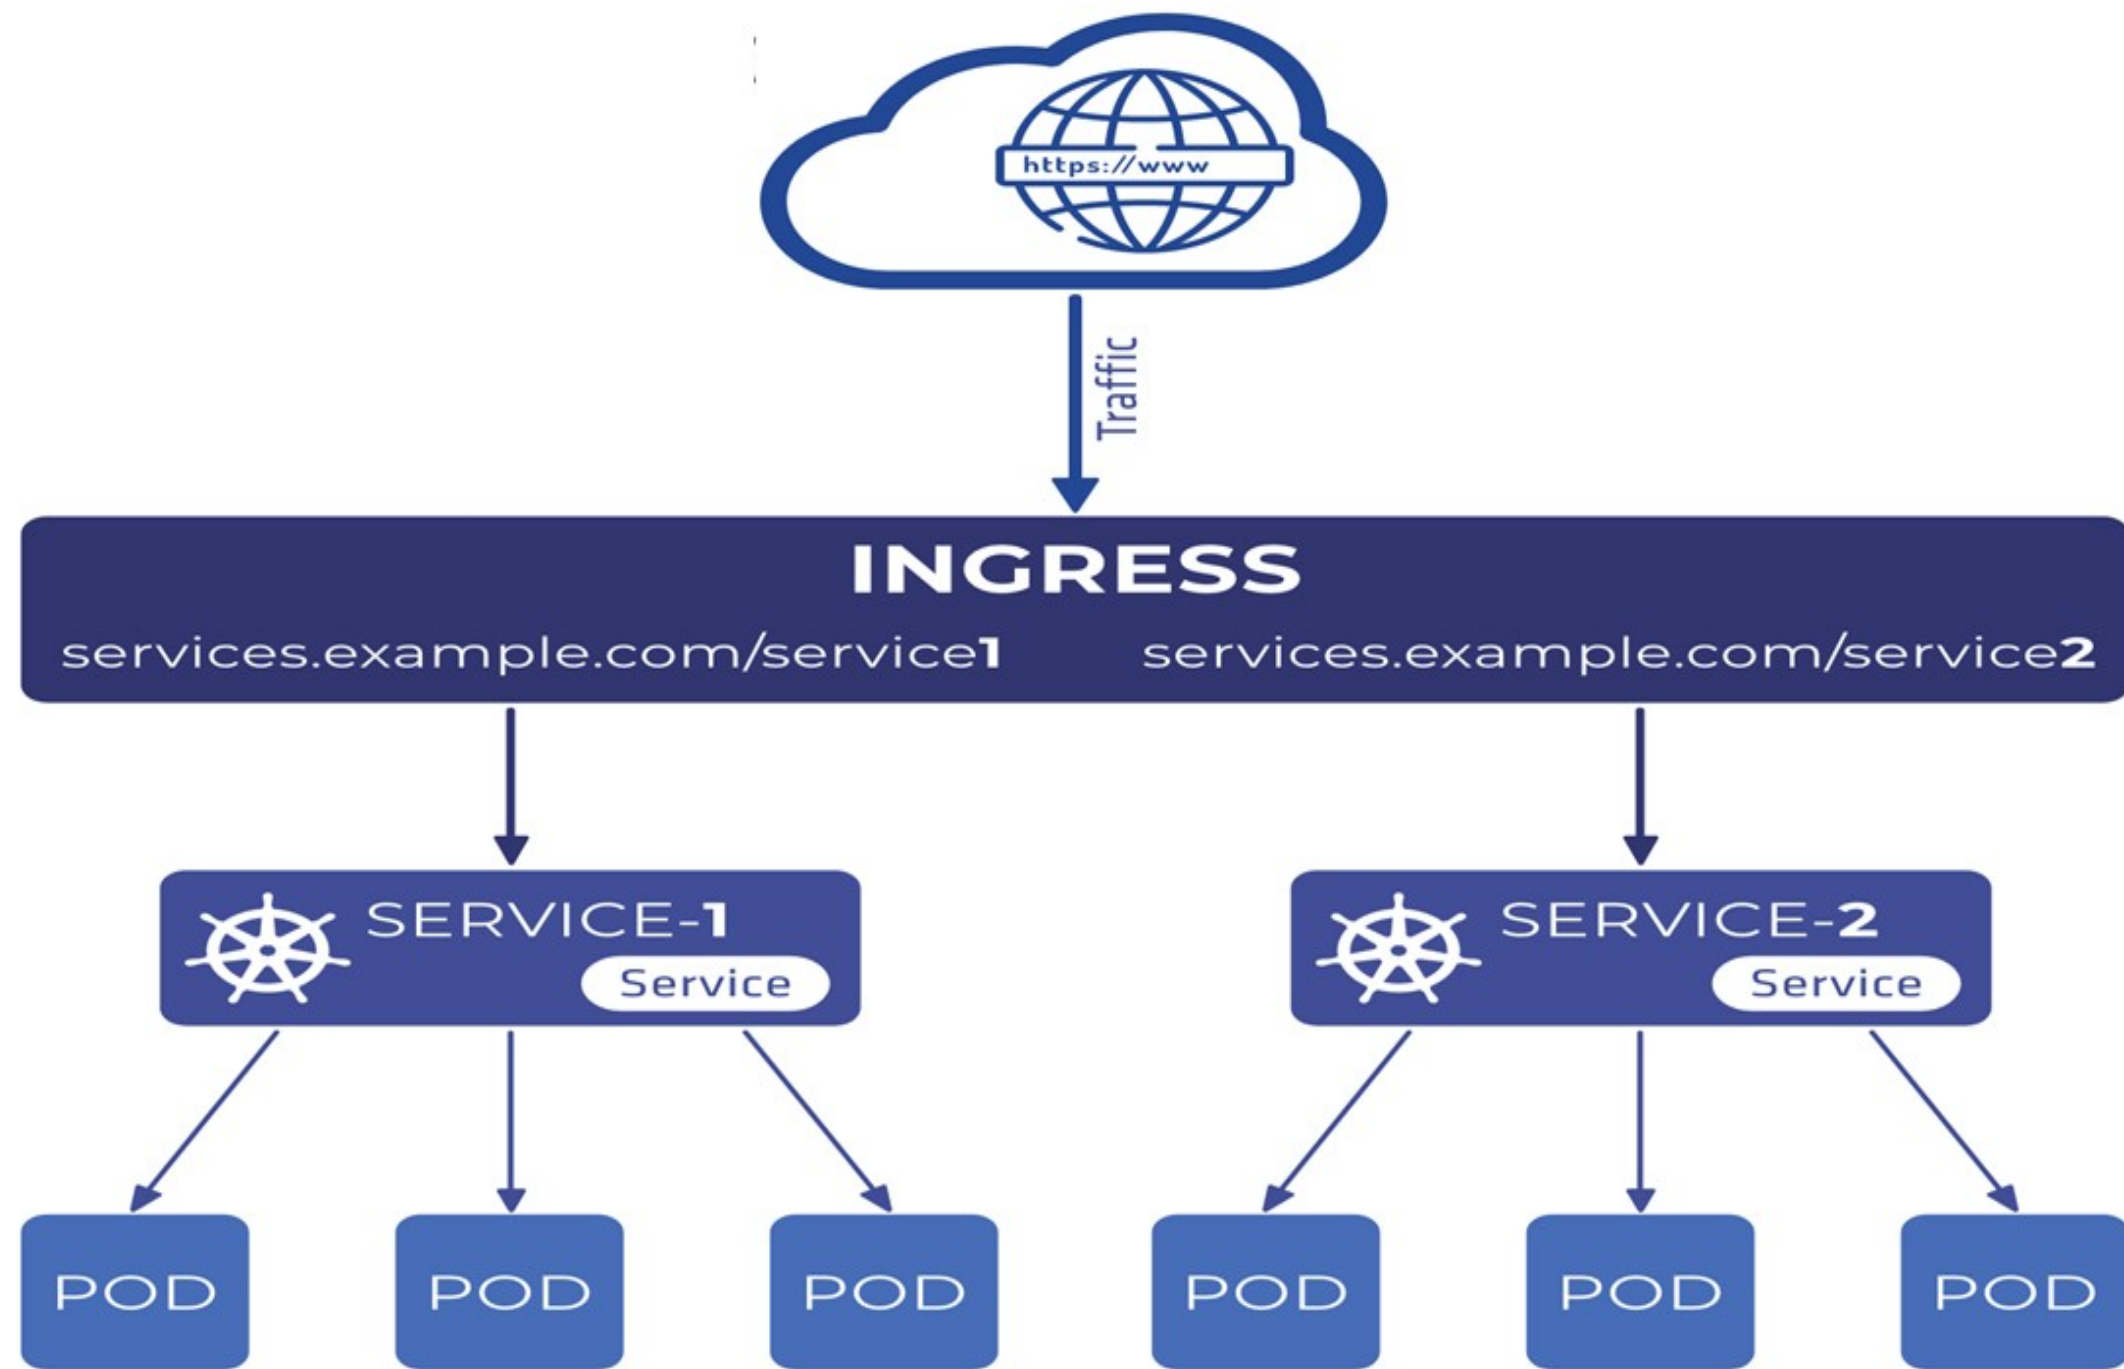
Loadbalancer service type automatically deploys an external loadbalancer for your service.

This loadbalancer service takes care of routing the requests to the underlying NodePort exposed.

Implementation of Load balancer depends upon your cloud service provider.

If using loadbalancer as a service on bare-metal, you have to provide own implementation of loadbalancer.





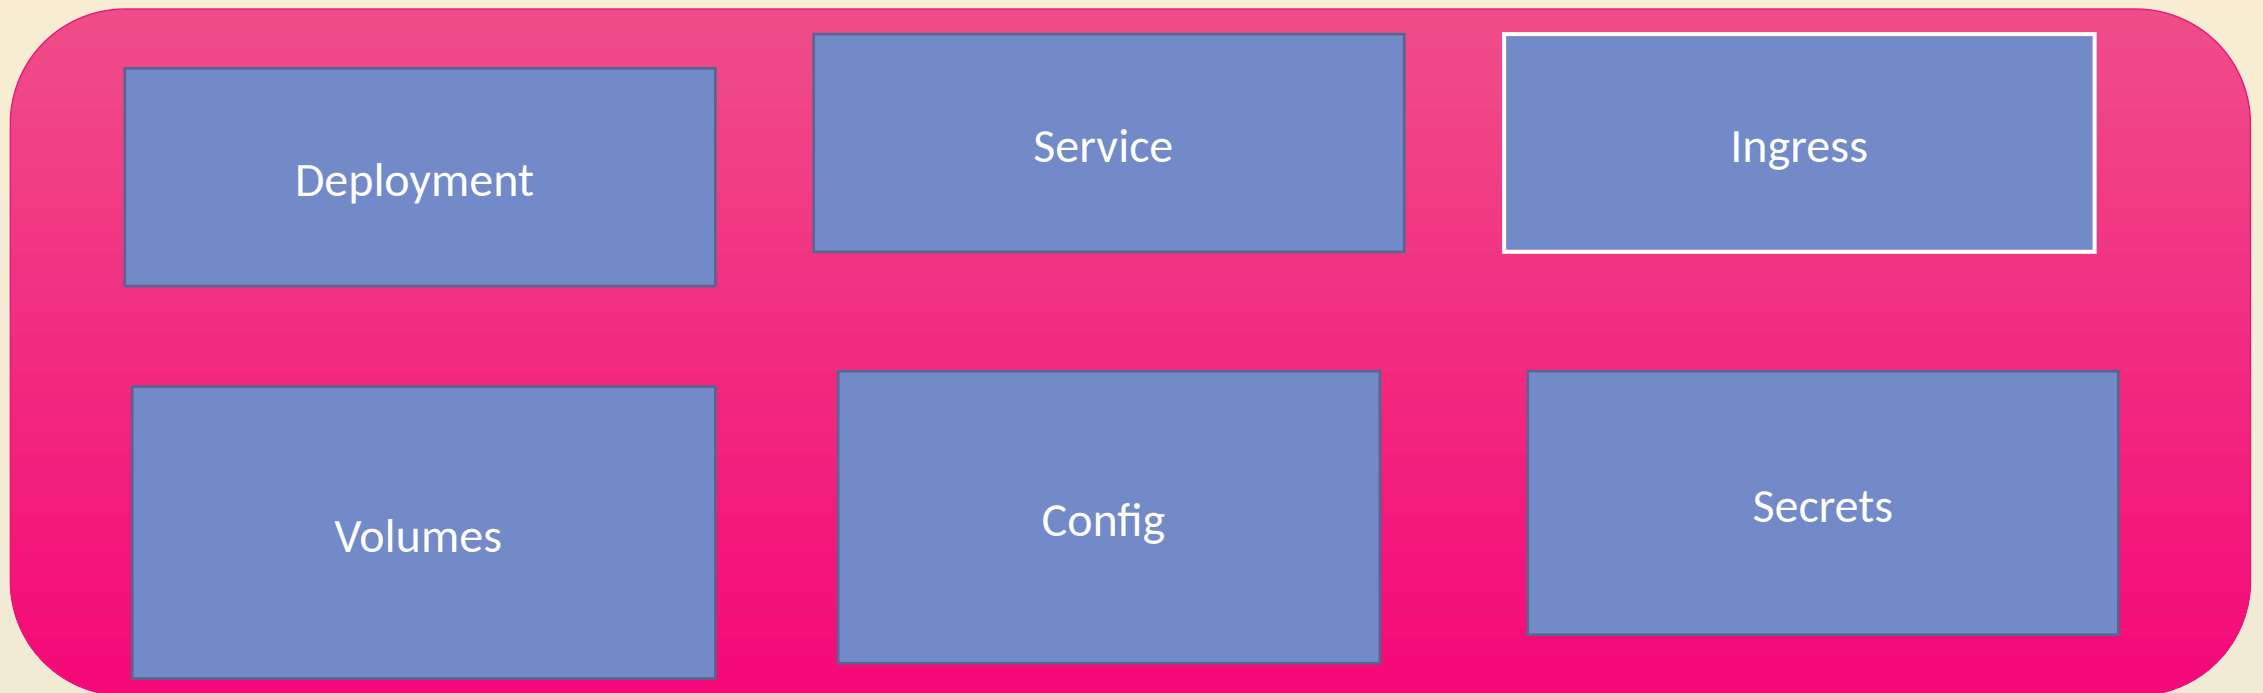
# HELM

Helm is a package manager for Kubernetes.

Helm is the K8s equivalent of yum or apt.

Helm deploys charts, which you can think of as a packaged application.

It is a collection of all your versioned, pre-configured application resources which can be deployed as one unit.

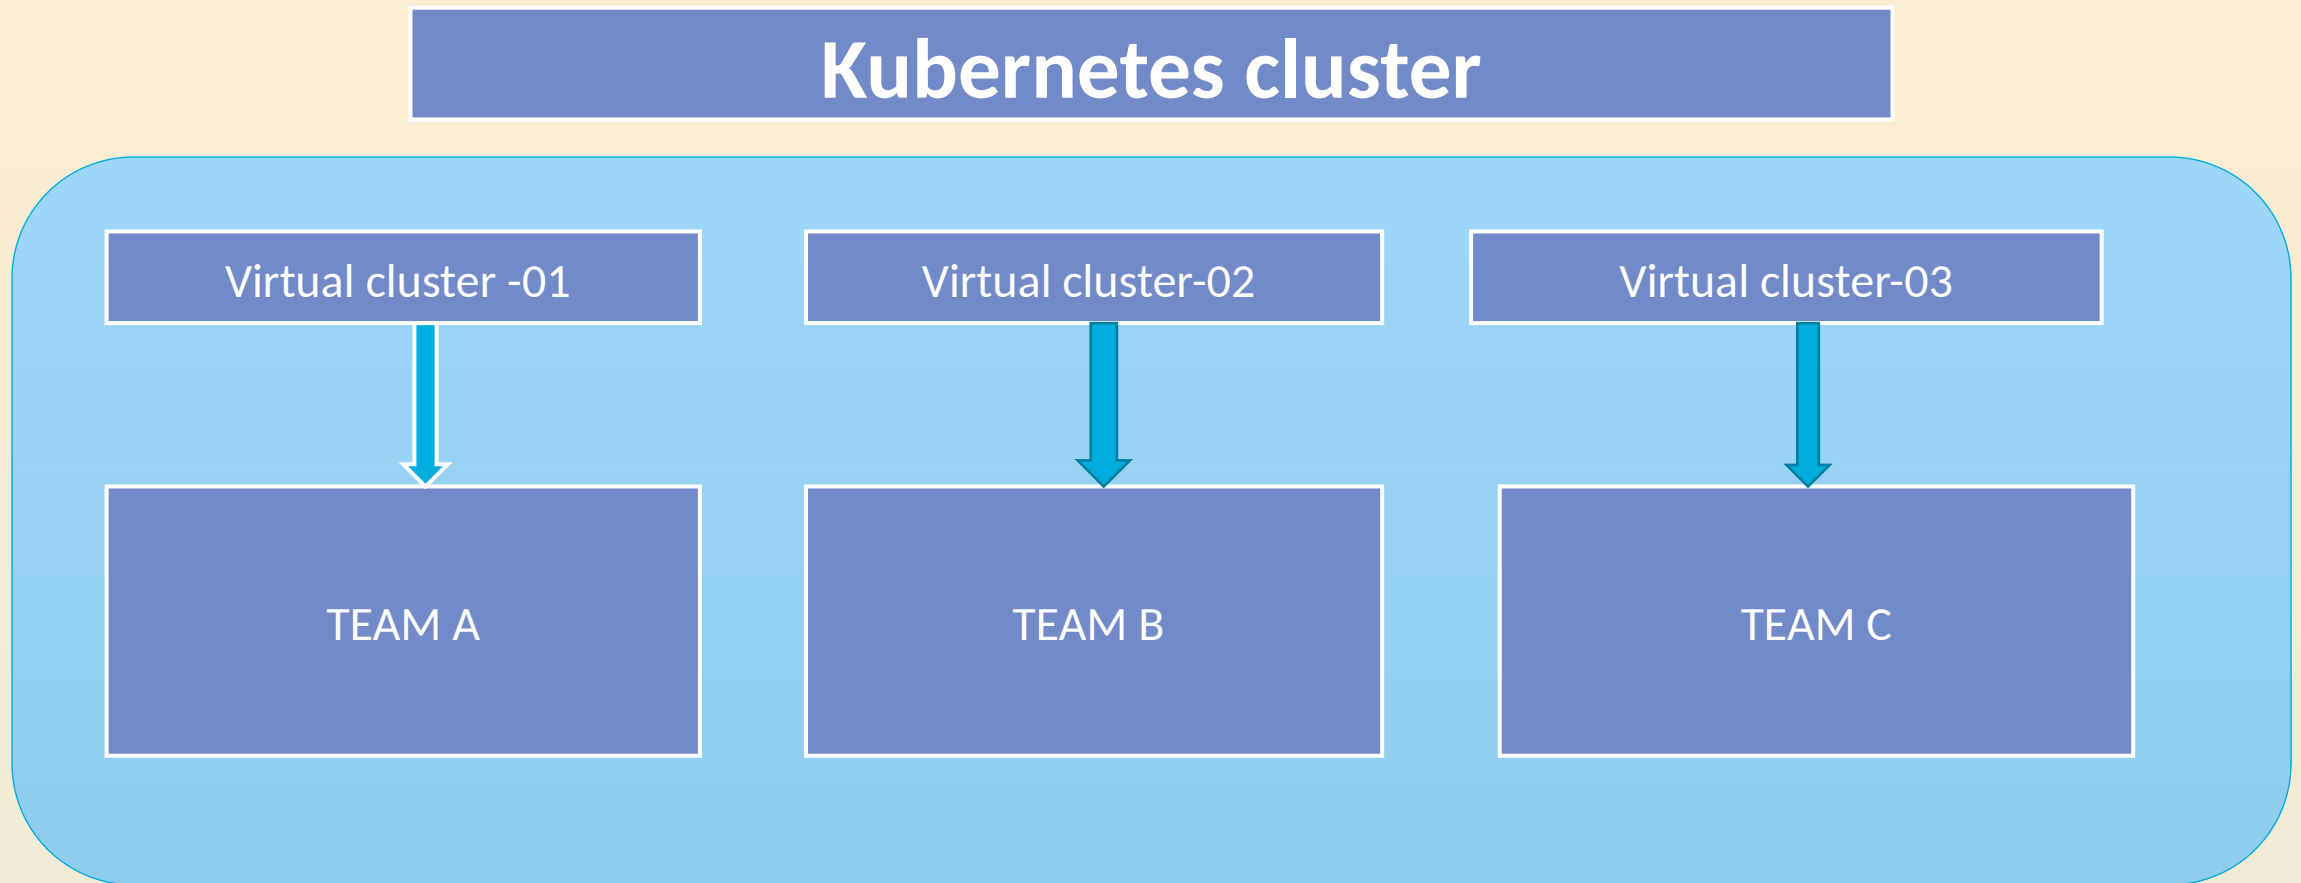




# KUBERNETES NAMESPACE

Kubernetes supports multiple virtual clusters for the same physical cluster.

These virtual clusters are called as Namespace.



**THANK YOU**